



Persistent Memory Programming

A Brief Tutorial

June 23, 2017

Andy Rudoff

Intel Corporation

DCG
Data Center Group



Links Used in This Tutorial

<http://pmem.io>

- Website for pmem programming, blogs, tutorials, examples

<https://github.com/pmem/nvml>

- Source for NVM Libraries supporting Windows, Linux in C and C++

<http://pmem.io/nvml/manpages/master/libpmemobj.3.html>

- libpmemobj man page (for C programming)

http://pmem.io/nvml/cpp_obj/master/cpp_html/index.html

- libpmemobj C++ interface documentation

<https://github.com/pmem/nvml/tree/master/src/examples>

- NVML examples, all buildable and runnable

<https://github.com/andyrudoff/pmemtutorial>

- The slides for this tutorial and the code examples (word frequency count)

Links to Additional Information

https://www.usenix.org/system/files/login/articles/login_summer17_07_rudoff.pdf

- An overview of persistent memory programming

<http://www.snia.org/PM>

SNIA Standards Portfolio

- NVM Programming Model v1.2a – Draft for public review
- NVM Programming Model v1.1- SNIA Technical Position
- NVM Programming Model v1.0 - SNIA Technical Position

SNIA Technical White Papers

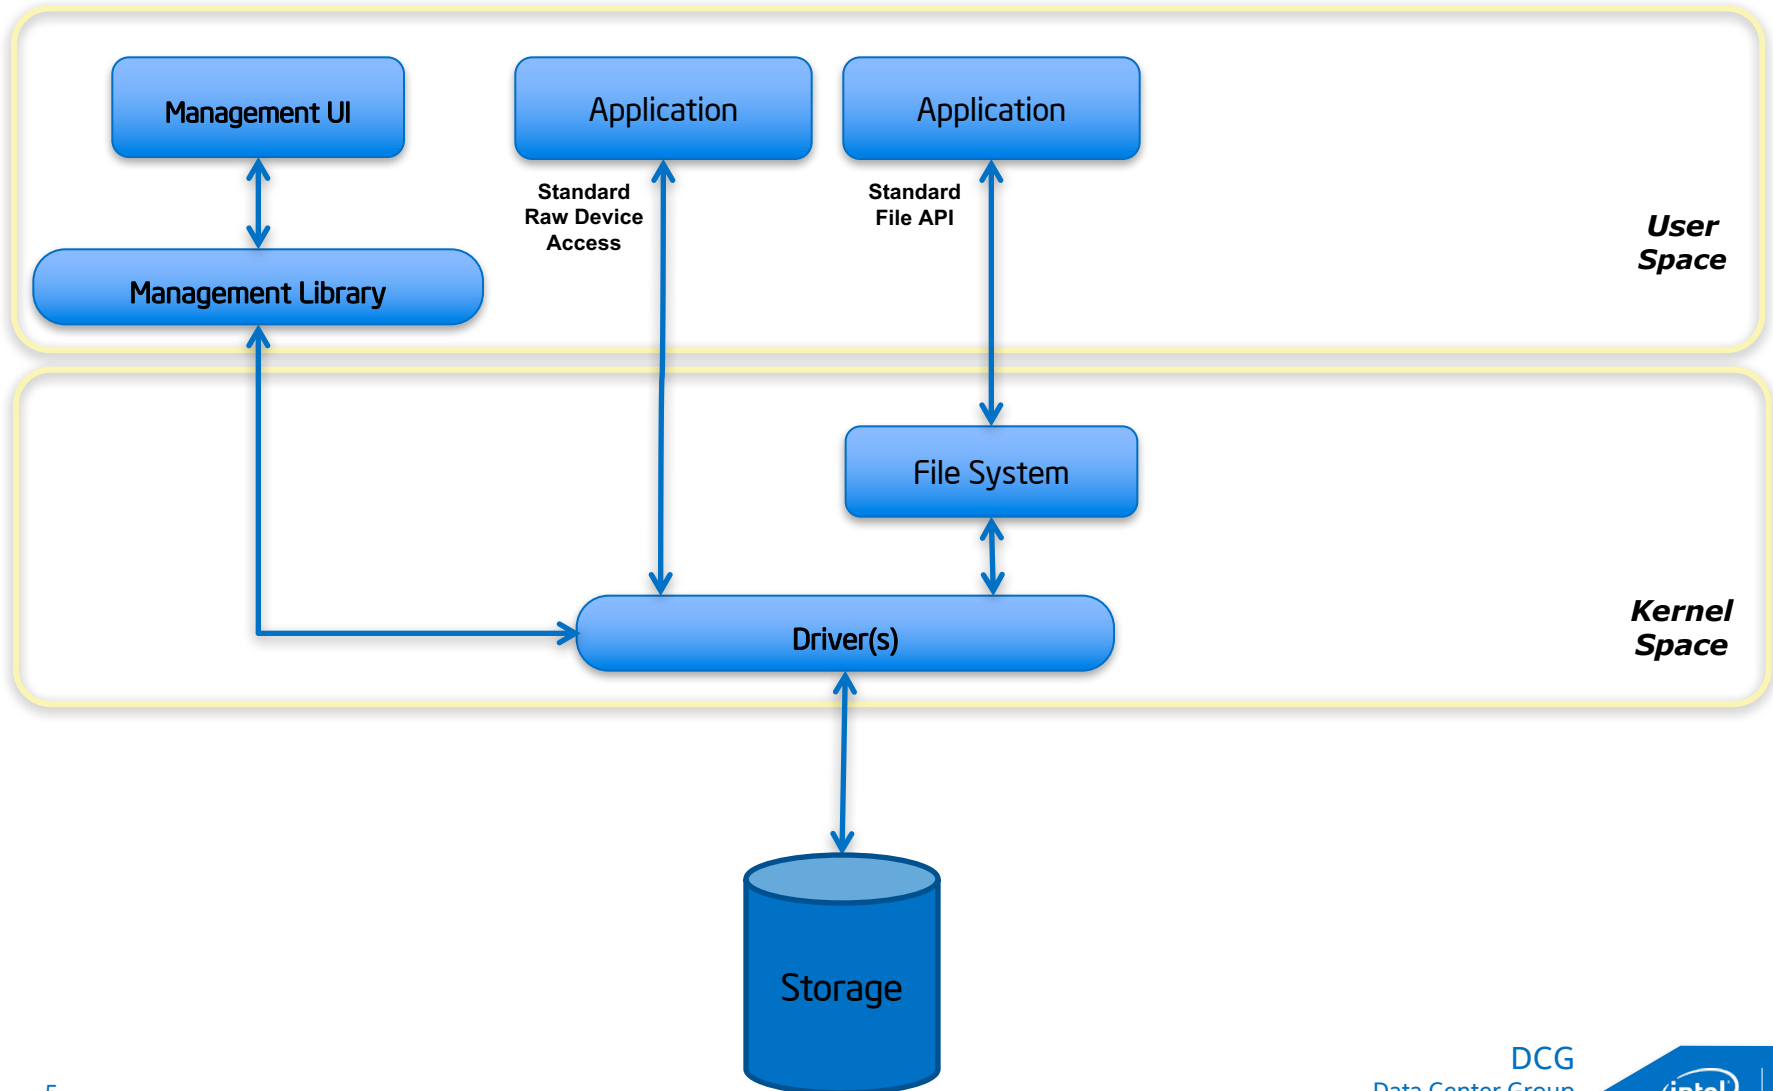
- NVM PM Remote Access for High Availability
- Persistent Memory Atomics and Transactions

SNIA Videos and Presentations

- The SNIA NVM Programming Model – Latest Developments and Challenges
- Persistent Memory Summit 2017

Background (abbreviated)

The Storage Stack (50,000ft view...)



A Programmer's View

(not just C programmers!)

```
fd = open("/my/file", O_RDWR);
```

```
...
```

```
count = read(fd, buf, bufsize);
```

```
...
```

```
count = write(fd, buf, bufsize);
```

```
...
```

```
close(fd);
```

"Buffer-Based"

A Programmer's View (mapped files)

```
fd = open("/my/file", O_RDWR);  
...  
base = mmap(NULL, filesize,  
            PROT_READ|PROT_WRITE, MAP_SHARED, fd, 0);  
close(fd);  
...  
base[100] = 'X';  
strcpy(base, "hello there");  
*structp = *base_structp;  
...
```

"Load/Store"

Memory-Mapped Files

What are memory-mapped files really?

- Direct access to the **page cache**
- Storage only supports block access (paging)

With load/store access, when does I/O happen?

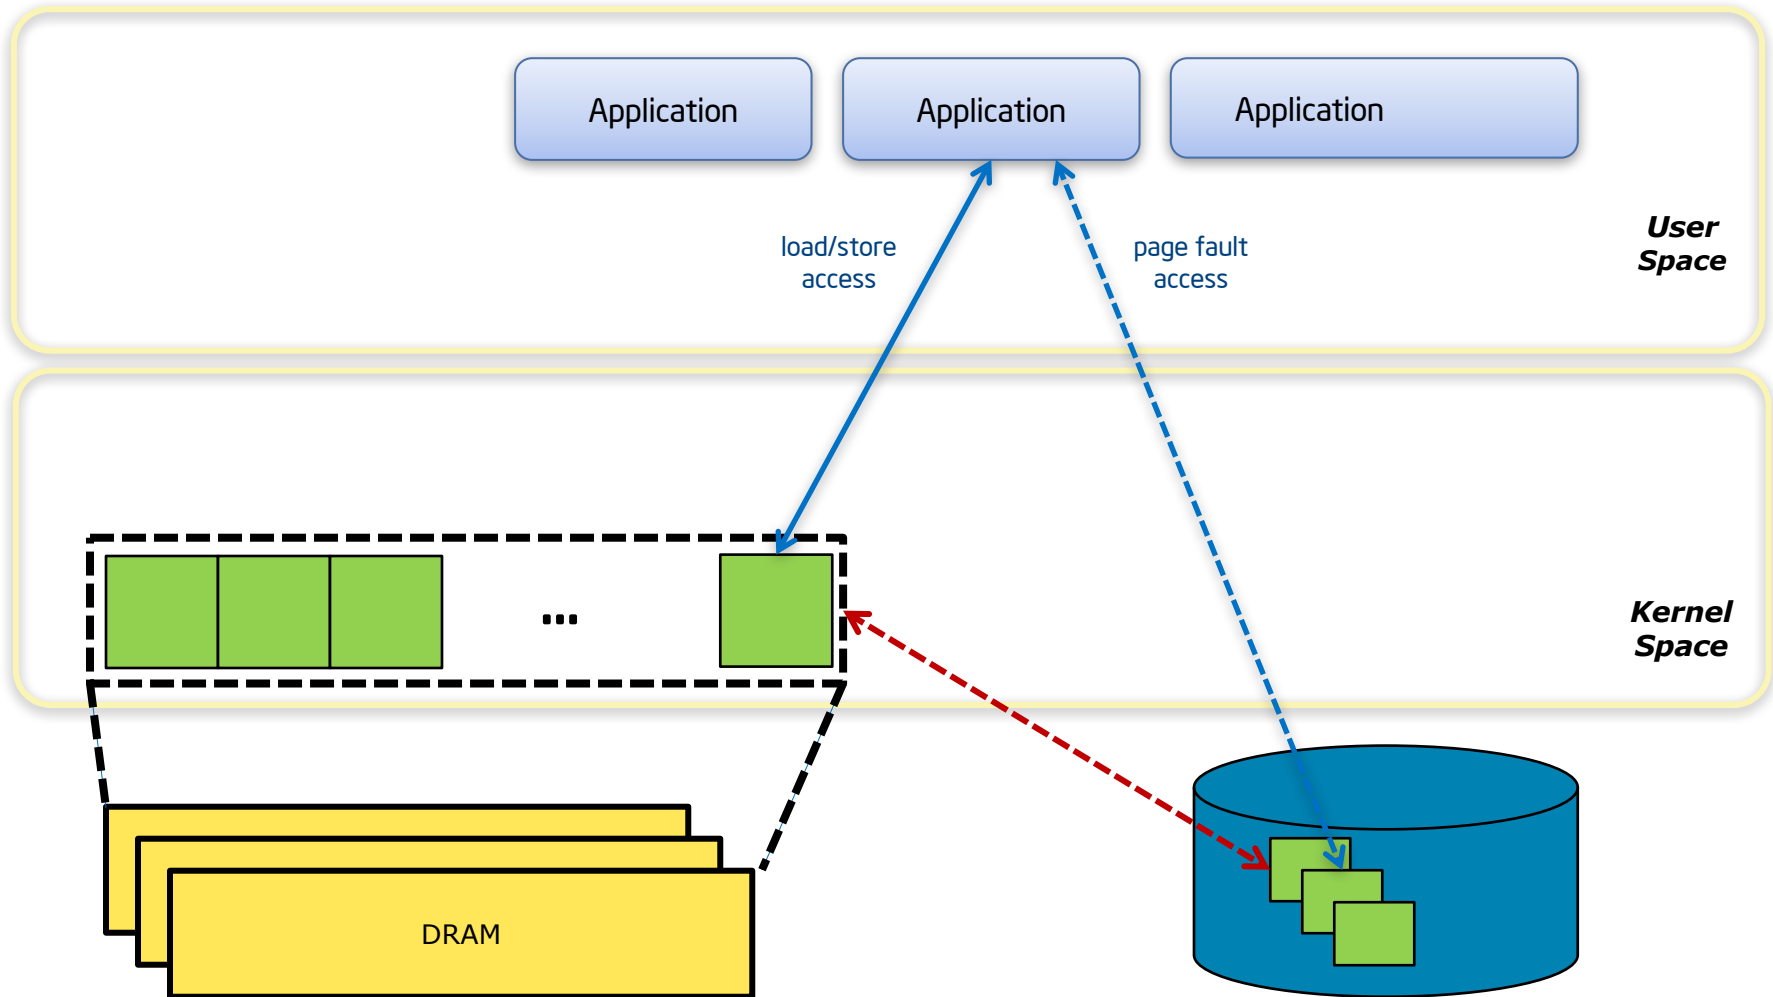
- Read faults/Write faults
- Flush to persistence

Not that commonly used or understood

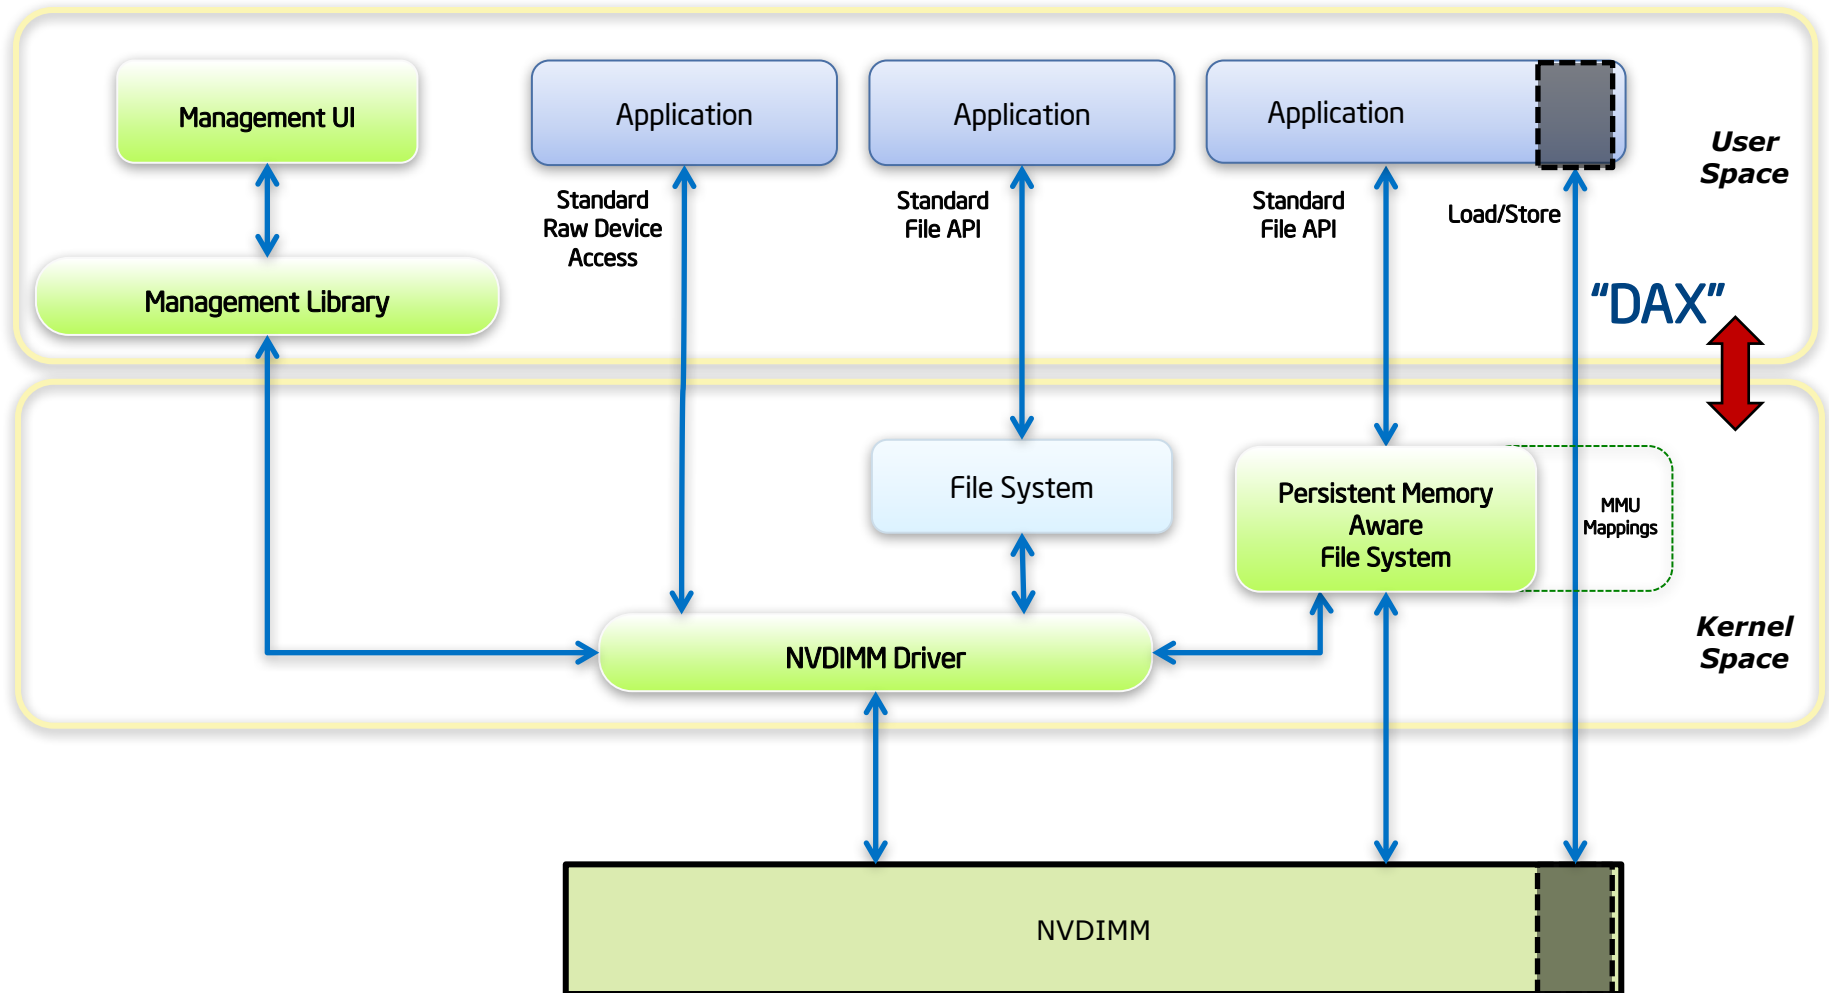
- Quite powerful
- Sometimes used without realizing it

Good reference: <http://nmmu.org/memory-faq.txt>

OS Paging

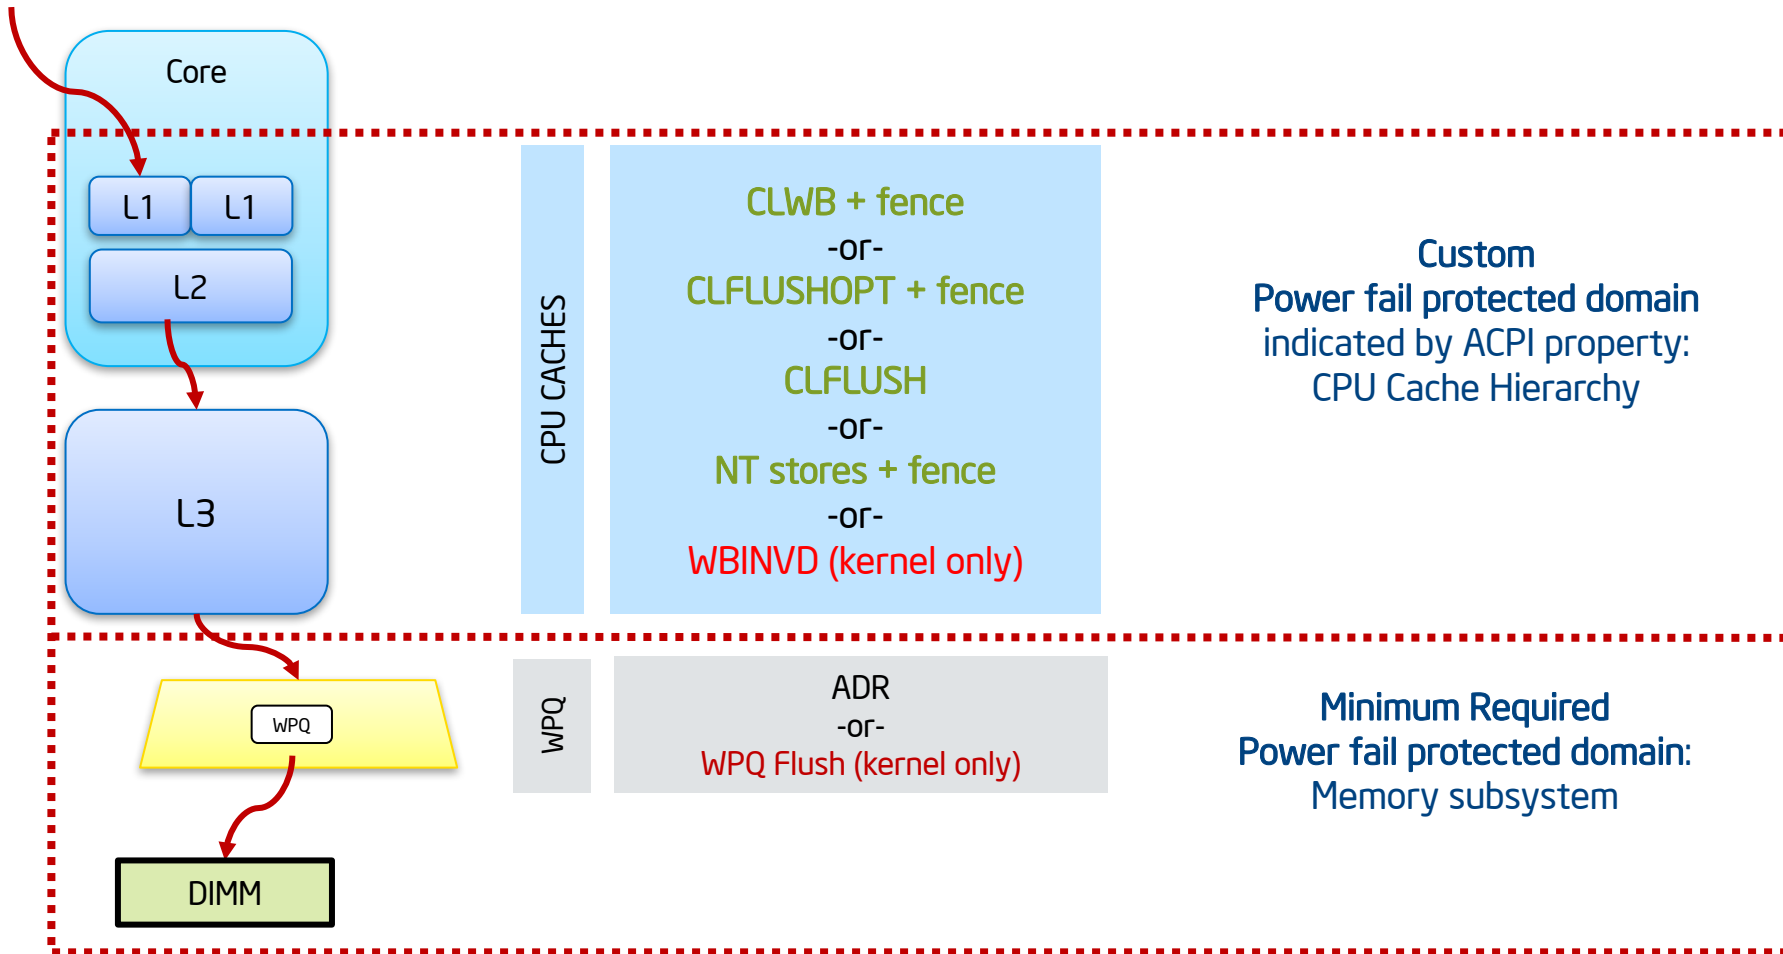


Exposing pmem to Applications



The Persistent Domain

MOV



Flushing for Application Programmers

Why is flushing required?

- Memory-mapped files have always worked this way:
 - Stores are not guaranteed persistent until flush API is called
 - Stores are *visible* before they are persistent

Do standard flushing APIs work with pmem?

- Yes, standard APIs work as expected
 - `msync()` on Linux
 - `FlushFileBuffers()` on Windows
 - The kernel will use instructions like CLWB as necessary

Can Applications just flush with CLWB from user space

- Only when supported by the kernel/file system
- Libraries like NVML determine when it is safe


State of Ecosystem Today

OS Detection of NVDIMMs	ACPI 6.0+
OS Exposes pmem to apps	DAX provides SNIA Programming Model Fully supported: <ul style="list-style-type: none">• Linux (ext4, XFS)• Windows (NTFS)
OS Supports Optimized Flush	Specified, but evolving (ask when safe) <ul style="list-style-type: none">• Linux: unsafe except Device DAX<ul style="list-style-type: none">• (and new file systems like NOVA)• Windows: safe
Remote Flush	Proposals under discussion (works today with extra round trip)
Deep Flush	Upcoming Specification
Transactions, Allocators	Built on above via libraries and languages: <ul style="list-style-type: none">• http://pmem.io Much more language support to do
Virtualization	All VMMs planning to support PM in guest (KVM changes upstream, Xen coming, others too...)

NVM Libraries

NVM Libraries: pmem.io

C/C++ on Linux and Windows

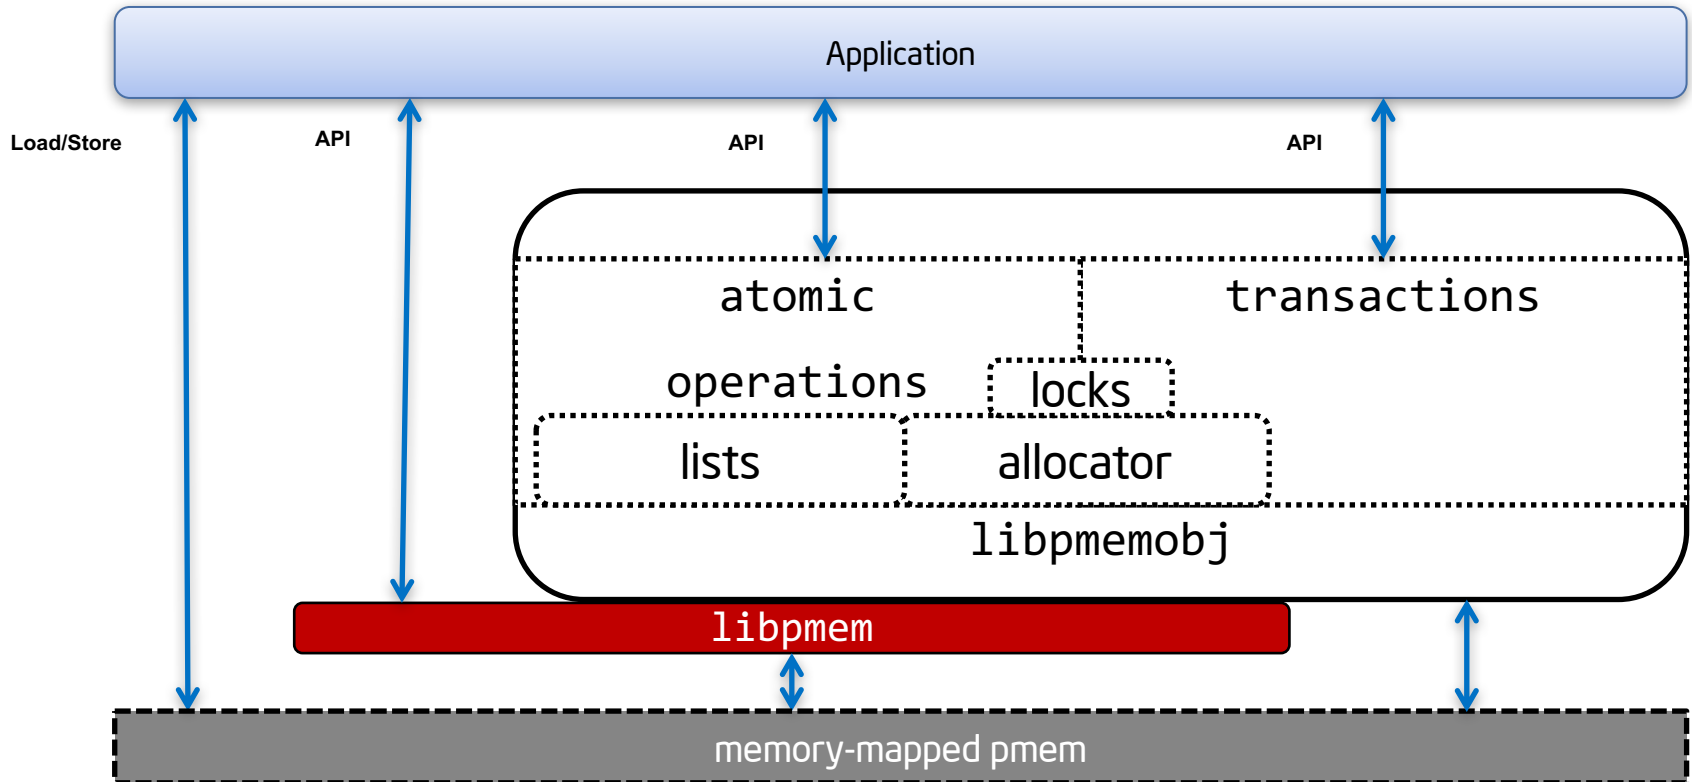


- Open Source
 - <http://pmem.io>
 - libpmem
 - libpmemobj
 - libpmemblk
 - libpmemlog
 - libvmem
- Transactional

More libraries being added to the suite over time

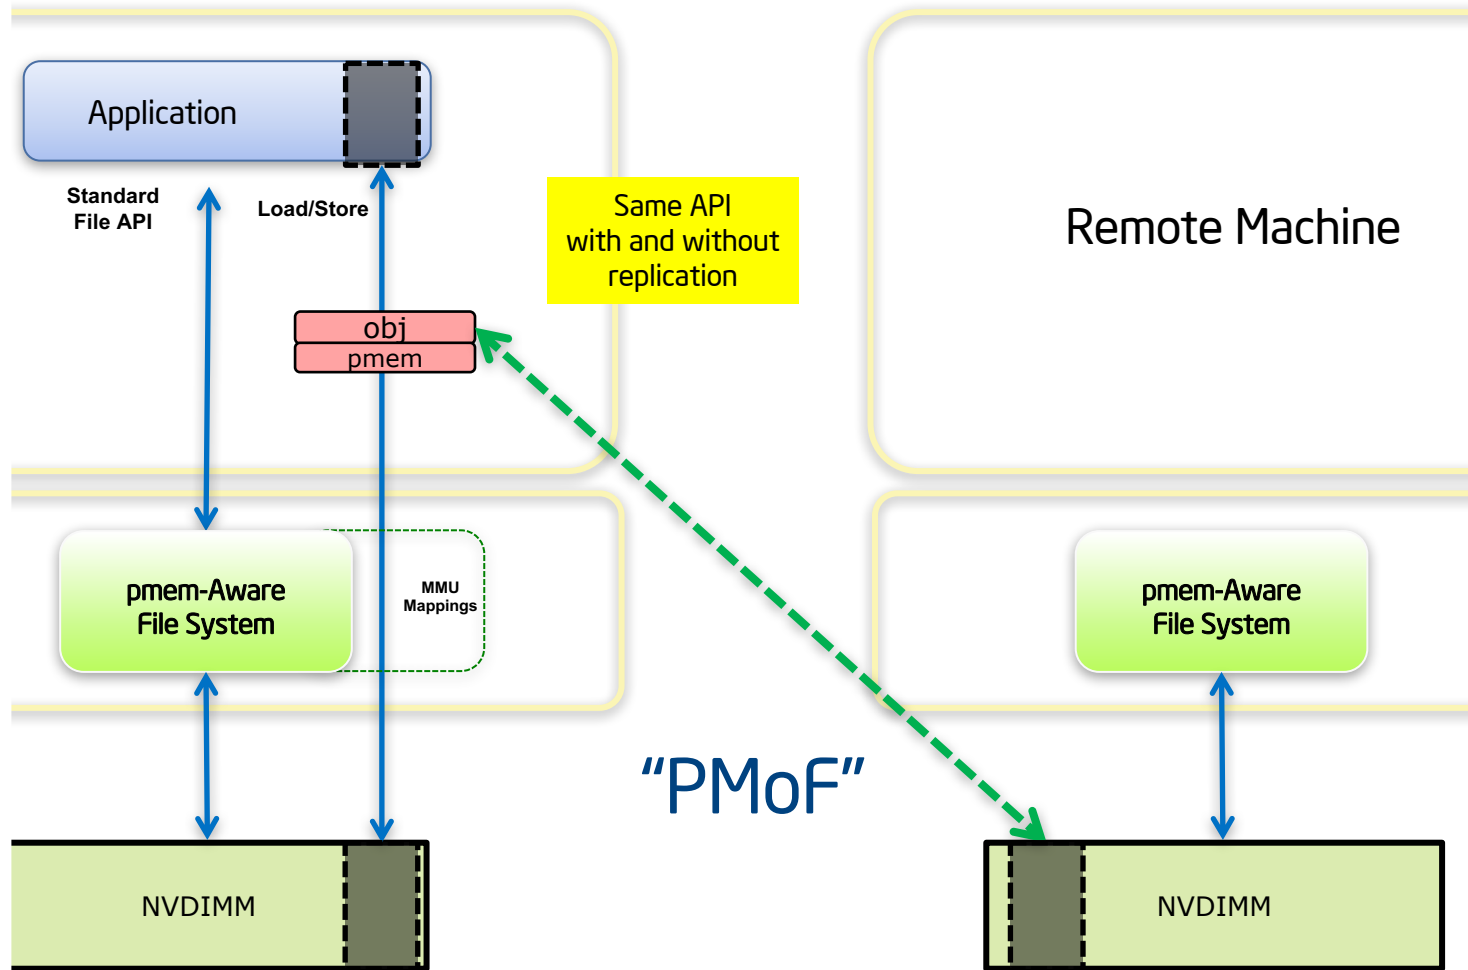
libpmemobj

“transactional object store”



Libpmemobj Replication: Application Transparent

(except for performance overhead)



Using libpmemobj Today

From C

- Fully validated, delivered on Linux, early access on Windows
- Can stick to pure API calls, but macros add:
 - Compile-time type safety
 - Transaction syntax, similar to try/catch

From C++

- Fully validated, delivered on Linux, early access on Windows
- Use C++ type system & syntax: much cleaner, less error-prone

From Java

- Persistent Containers for Java (Experimental)

From Python

- PyNVM (Experimental)

Other work

- valgrind (and a similar tool coming from Intel)
- JavaScript (Pre-release)

Emulating Persistent Memory

The programming model builds on memory-mapped files

- So development on memory-mapped files works fine
 - NVML will use `msync()` to flush to persistence
 - Non-optimal performance
- Use any 64-bit Linux or Windows

For benchmarking:

- <http://pmem.io/2016/02/22/pm-emulation.html>
- Distros like Fedora 24 are built with DAX/pmem
 - Avoids making you build a kernel
 - Also avoids building NVML

Using NVML on Fedora 24 or later...

```
fedora24 # dnf install libpmemobj-devel
```

```
Last metadata expiration check: 0:08:18 ago on Wed Sep 14 14:58:49 2016.
```

```
Dependencies resolved.
```

=====				
Package	Arch	Version	Repository	Size
=====				
Installing:				
libpmem	x86_64	1.1-1.fc24	updates	29 k
libpmem-devel	x86_64	1.1-1.fc24	updates	43 k
libpmemobj	x86_64	1.1-1.fc24	updates	66 k
libpmemobj-devel	x86_64	1.1-1.fc24	updates	112 k

```
Transaction Summary
```

```
=====
```

```
Install 4 Packages
```

```
Total download size: 251 k
```

```
Installed size: 527 k
```

```
Is this ok [y/N]: y
```

Downloading Packages:

(1/4): libpmem-devel-1.1-1.fc24.x86_64.rpm	81 kB/s		43 kB	00:00
(2/4): libpmemobj-devel-1.1-1.fc24.x86_64.rpm	184 kB/s		112 kB	00:00
(3/4): libpmem-1.1-1.fc24.x86_64.rpm	209 kB/s		29 kB	00:00
(4/4): libpmemobj-1.1-1.fc24.x86_64.rpm	98 kB/s		66 kB	00:00

Total	153 kB/s		251 kB	00:01
-------	----------	--	--------	-------

Running transaction check

Transaction check succeeded.

Running transaction test

Transaction test succeeded.

Running transaction

Installing	: libpmem-1.1-1.fc24.x86_64	1/4
Installing	: libpmem-devel-1.1-1.fc24.x86_64	2/4
Installing	: libpmemobj-1.1-1.fc24.x86_64	3/4
Installing	: libpmemobj-devel-1.1-1.fc24.x86_64	4/4
Verifying	: libpmemobj-devel-1.1-1.fc24.x86_64	1/4
Verifying	: libpmem-devel-1.1-1.fc24.x86_64	2/4
Verifying	: libpmemobj-1.1-1.fc24.x86_64	3/4
Verifying	: libpmem-1.1-1.fc24.x86_64	4/4

Installed:

libpmem.x86_64 1.1-1.fc24	libpmem-devel.x86_64 1.1-1.fc24
libpmemobj.x86_64 1.1-1.fc24	libpmemobj-devel.x86_64 1.1-1.fc24

Complete!

The `pmempool` command

(`nvml-tools` Package)

`pmempool-info(1)`

Prints information and statistics in human-readable format about specified pool.

`pmempool-check(1)`

Checks pool's consistency and repairs pool if it is not consistent.

`pmempool-create(1)`

Creates a pool of specified type with additional properties specific for this type of pool.

`pmempool-dump(1)`

Dumps usable data from pool in hexadecimal or binary format.

`pmempool-rm(1)`

Removes pool file or all pool files listed in poolset configuration file.

`pmempool-convert(1)`

Updates the pool to the latest available layout version.

Tour Through <http://pmem.io> and NVML Source Tree

Essential libpmem Knowledge

libpmem examples

Source: <https://github.com/pmem/nvml/tree/master/src/examples/libpmem>

```
/*
 * simple_copy.c -- show how to use pmem_memcpy_persist()
 *
 * usage: simple_copy src-file dst-file
 *
 * Reads 4k from src-file and writes it to dst-file.
 */

/* create a pmem file and memory map it */
if ((pmemaddr = pmem_map_file(argv[2], BUF_LEN,
                             PMEM_FILE_CREATE|PMEM_FILE_EXCL,
                             0666, &mapped_len, &is_pmem)) == NULL) {
    perror("pmem_map_file");
    exit(1);
}
```

Using `is_pmem`

```
if (is_pmem) {  
    pmem_memcpy_persist(pmemaddr, buf, cc);  
} else {  
    memcpy(pmemaddr, buf, cc);  
    pmem_msync(pmemaddr, cc);  
}
```


libvmem Example

Volatile use of Persistent Memory

```
if ((vmp = vmem_create("/pmem-fs", VMEM_MIN_POOL)) == NULL) {  
    perror("vmem_create");  
    exit(1);  
}  
  
if ((ptr = vmem_malloc(vmp, 100)) == NULL) {  
    perror("vmem_malloc");  
    exit(1);  
}  
  
strcpy(ptr, "hello, world");  
  
/* give the memory back */  
vmem_free(vmp, ptr);  
  
/* ... */  
  
vmem_delete(vmp);
```

libpmemobj Examples

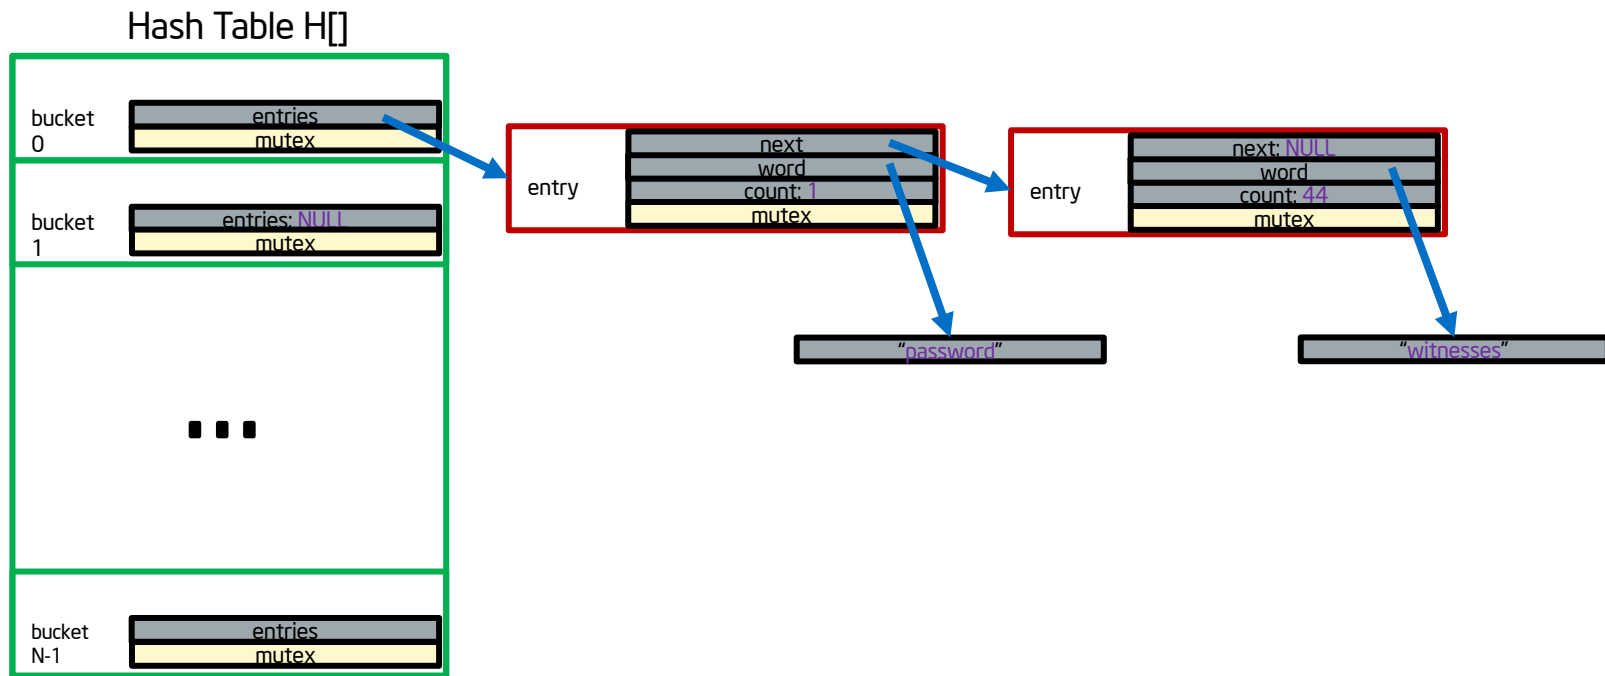
Simple C program to build example on (nothing related to pmem yet)



freq.c

```
$ freq -p words.txt
1 is
1 all
1 for
2 to
1 men
1 good
2 the
1 come
1 their
1 Now
1 time
1 country
1 aid
1 of
```


Adding multi-threading support (nothing related to pmem yet)



freq_mt.c

```
$ freq_mt -p words.txt words.txt words.txt
3 is
3 all
3 for
6 to
3 men
3 good
6 the
3 come
3 their
3 Now
3 time
3 country
3 aid
3 of
```


The *Root Object*



root object:


- assume it is always there
- created first time accessed
- initially zeroed

Using the Root Object

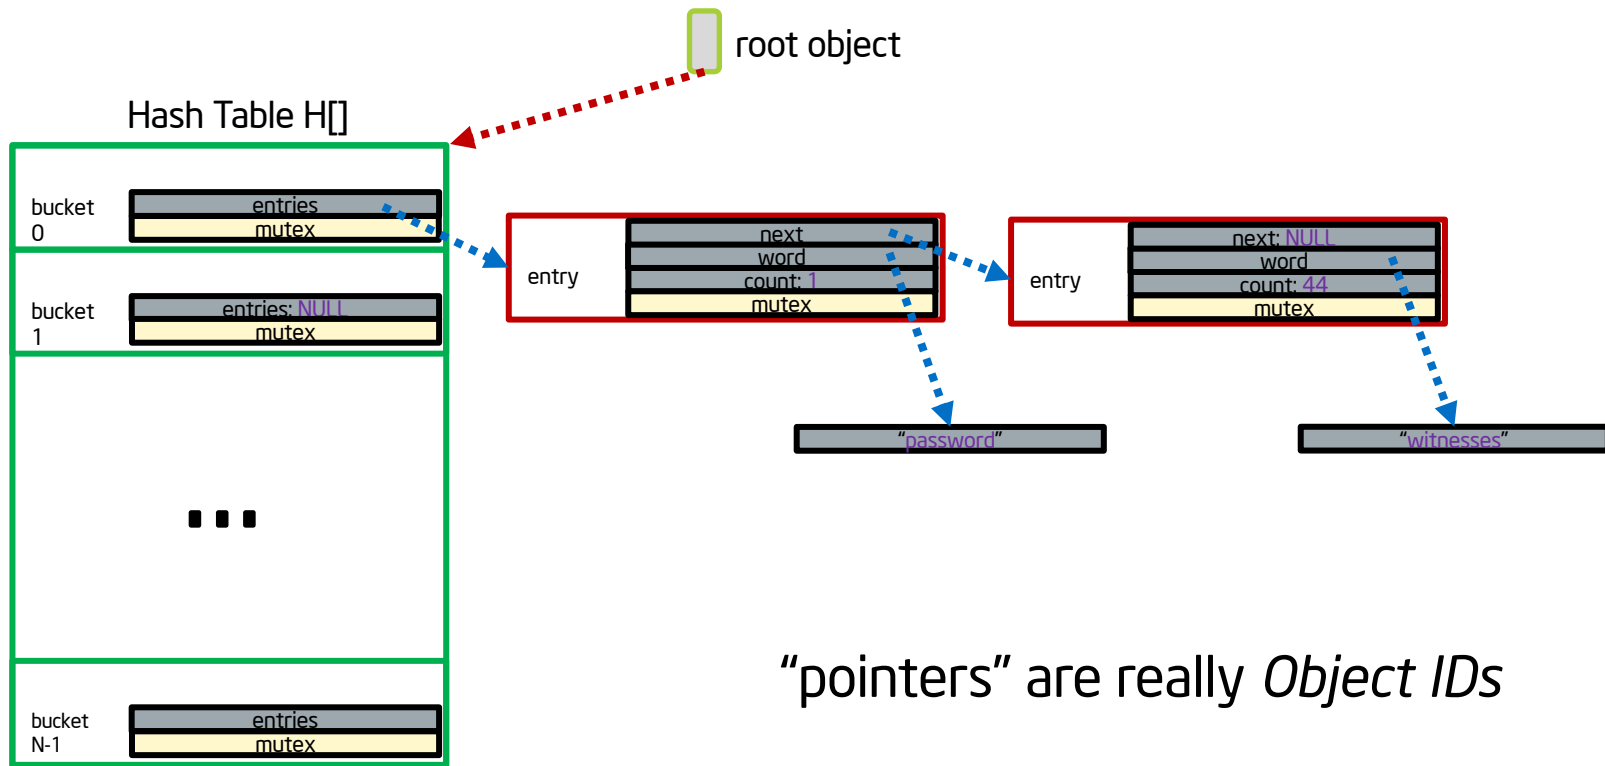


Link pmem data structures in pool
off the root object to find
them on each program run

“pointers” are really *Object IDs*



Moving data the example to pmem



"pointers" are really *Object IDs*

C Programming with libpmemobj

Transaction Syntax

```
TX_BEGIN(Pop) {  
    /* the actual transaction code goes here... */  
} TX_ONCOMMIT {  
    /*  
     * optional - executed only if the above block  
     * successfully completes  
     */  
} TX_ONABORT {  
    /*  
     * optional - executed if starting the transaction fails  
     * or if transaction is aborted by an error or a call to  
     * pmemobj_tx_abort()  
     */  
} TX_FINALLY {  
    /*  
     * optional - if exists, it is executed after  
     * TX_ONCOMMIT or TX_ONABORT block  
     */  
} TX_END /* mandatory */
```

Properties of Transactions

Powerfail
Atomicity

Multi-Thread
Atomicity

```
TX_BEGIN_PARAM(Pop, TX_PARAM_MUTEX, &D_RW(ep)->mtx, TX_PARAM_NONE) {  
    TX_ADD(ep);  
    D_RW(ep)->count++;  
} TX_END
```

Caller must
instrument code
for undo logging

Persistent Memory Locks

- Want locks to live near the data they protect (i.e. inside structs)
- Does the state of locks get stored persistently?
 - Would have to flush to persistence when used
 - Would have to recover locked locks on start-up
 - Might be a different program accessing the file
 - Would run at pmem speeds
- PMEMmutex
 - Runs at DRAM speeds
 - Automatically initialized on pool open

freq_pmem.c

```
$ pmempool create obj --layout=freq -s 1G freqcount  
  
$ freq_pmem_print freqcount  
  
$ freq_pmem freqcount words.txt words.txt words.txt  
  
$ freq_pmem_print freqcount  
3 is  
3 all  
3 for  
6 to  
3 men  
3 good  
6 the  
...
```

C++ Programming with libpmemobj

C++ Queue Example: Declarations

```
/* entry in the queue */  
struct pmem_entry {  
    persistent_ptr<pmem_entry> next;  
    p<uint64_t> value;  
};
```

`persistent_ptr<T>`

Pointer is really a position-independent
Object ID in pmem.

Gets rid of need to use C macros like `D_RW()`

`p<T>`

Field is pmem-resident and needs to be
maintained persistently.

Gets rid of need to use C macros like `TX_ADD()`

C++ Queue Example: Transaction

```
void push(pool_base &pop, uint64_t value) {  
    transaction::exec_tx(pop, [&] {  
        auto n = make_persistent<pmem_entry>();  
  
        n->value = value;  
        n->next = nullptr;  
        if (head == nullptr) {  
            head = tail = n;  
        } else {  
            tail->next = n;  
            tail = n;  
        }  
    });  
}
```

Transactional
(including allocations & frees)

freq_pmem_cpp.c

```
$ freq_pmem_cpp freqcount words.txt words.txt words.txt
```

```
$ freq_pmem_print freqcount
```

```
6 is
```

```
6 all
```

```
6 for
```

```
12 to
```

```
6 men
```

```
6 good
```

```
12 the
```

```
6 come
```

```
6 their
```

```
6 Now
```

```
6 time
```

```
6 country
```

```
6 aid
```

```
6 of
```

Future C++ Programming with libpmemobj

Persistent memory containers

A proof of concept under way.

- Targeting libc++ and libstdc++

```
329     typedef std::vector<foo, nvml::obj::allocator<foo>> pvector;
330
331     struct root {
332         persistent_ptr<pvector> my_vector;
333     };
334
335     nvobj::pool<root> pop = nvobj::pool<root>::open(path, "layout");
336
337     transaction::exec_tx(pop, [&] {
338         auto root = pop.get_root();
339
340         root->my_vector->emplace_back(0xDEADBEEF);
341         root->my_vector->push_back(foo(0xBADA55));
342
343         for(auto el : root->my_vector)
344             std::cout << el << std::endl;
345     });
```

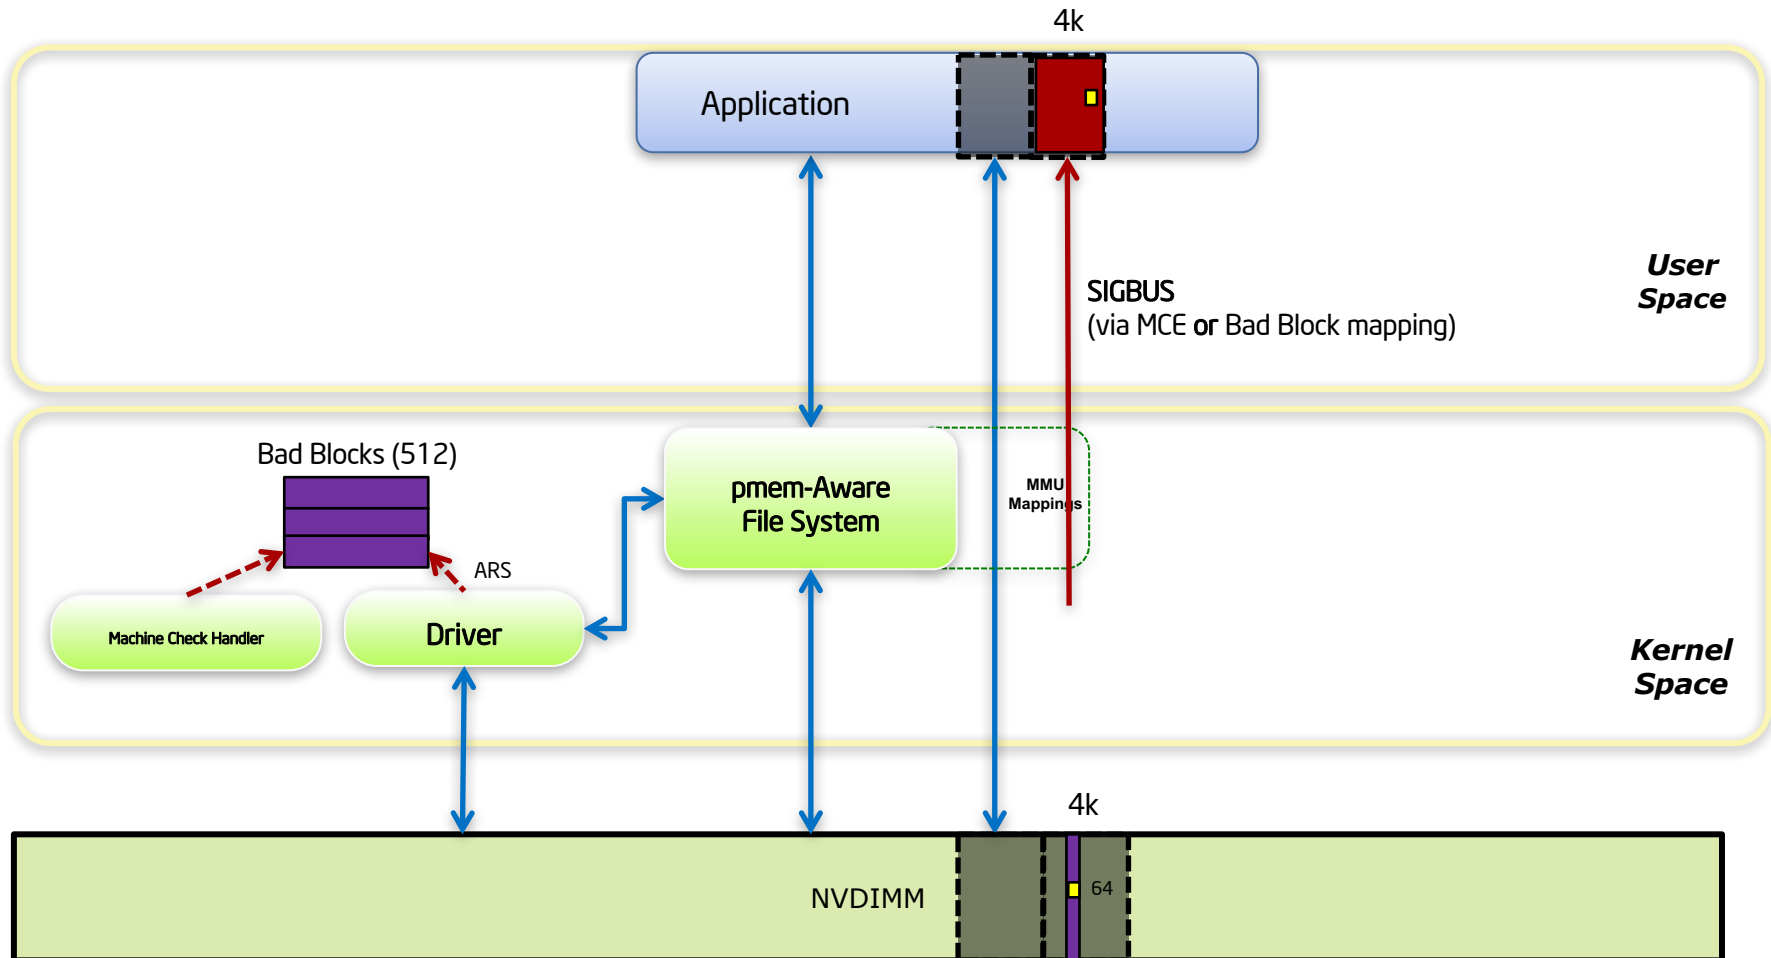
The Inconvenient Truth

Flush On Fail Fail

- ADR Failure Detection
 - Once detected, what SW knows the recovery action?
 - Usually the App
 - Ordering information lost
 - libpmem additions in progress
- pmem-Based Block Storage Errors
 - Without extra hardware, comes in as Machine Check
 - With hardware, can return to driver model
- Uncorrectables...

Linux Example:

Blast radius 64B → 4k



Summary

Tutorial Summary

Persistent Memory

- Emerging technology, game changing, large capacity
- New programming models allow greater leverage

NVM Libraries

- <http://pmem.io>
- Convenience, not a requirement
- Transactions, memory allocation, language support
- More coming

We don't know all the answers yet

- The next few years are going to be pretty exciting!

